

MongoDB

Learning MongoDB

PRAVEEN NAIR

Introduction to MongoDB

MongoDB is a document database.

MongoDB is a non-relational, non-tabular database.

Relational data is stored differently.

Instead of having multiple tables all the related data are stored together.

In MongoDB, tables are called collections.

MongoDB can be installed locally or in cloud called MongoDB Atlas

Mongosh or Compass can be used to query MongoDB

Advantages of MongoDB

Flexibility: MongoDB is schema-less, meaning you don't need to design a schema for the database.

Scalability: MongoDB can be horizontally scaled by distributing data across multiple servers, a process called sharding.

Performance: MongoDB is fast at inserting or updating large numbers of records. It also supports geospatial efficiently.

MongoDB Community Server

<https://www.mongodb.com/try/download/community>

Choose MSI

MongoDB Shell Download

<https://www.mongodb.com/try/download/shell>

Choose MSI

Connect to local mongodb

Type mongosh –version

Type mongosh to get prompt

show dbs

use myproj to create or access new db

db.dropDatabase("dbname") to delete database (or db.dropDatabase())

show collections

db.createCollection("employees")

db.employees.drop() to delete collection

db.restaurant.renameCollection('restaurants') //rename collection

Inserting Data

```
db.employees.insertOne({  
    name: "John Smith",  
    email: "john@gmail.com",  
    department: "IT",  
    salary: 1456,  
    location: ["FL", "OH"],  
    date: Date()  
})  
  
db.employees.find()
```

Inserting Multiple Data

```
db.employees.insertMany([{
  name: "Mike Joseph",
  email: "mike@gmail.com",
  department: "IT",
  salary: 2456,
  location: ["FL", "TX"],
  date: Date()
},
{ name: "Cathy G",
  email: "cathy@gmail.com",
  department: "IT",
  salary: 3456,
  location: ["AZ", "TX"],
  date: Date()
}])
```

Find Data – part 1

```
db.employees.find() //returns first 20, then type it for more documents  
db.employees.find().skip(2)  
db.employees.findOne()  
db.users.find().sort({name:1}) //sorting -1 for reverse  
db.users.find().limit(1) //returns 1 document sort by object id  
db.users.find().sort({name:1}).limit(3)  
db.employees.find( {department: "IT"} )  
db.users.find({name:"Cathy",pass:"1234"}) //two condition  
db.employees.find({}, {_id: 0, salary: 1, date: 1}) //cannot give 0  
db.users.find({},{_id:false,name:true}) //cannot give false  
db.employees.find({}, {_id: 0, salary: 0, date: 1}) //either use 0 or 1, can't use both
```

Find Data – part 2

```
db.users.find({'address.city':'Gwenborough'}) //query nested documents
db.users.find({address.geo.lat:-37.3159})
db.employees.find({'location':'TX'}) //where location : ['FL','TX']
db.users.find().count()
db.employees.find({}, {"dept": "$department", email:1,salary:1}) //dept is alias
db.users.find({'address.city':'Gwenborough'}) //query nested documents
db.users.find({address.geo.lat:-37.3159})
db.employees.find({'location':'TX'}) //where location : ['FL','TX']
```

returns first 20, then type it for more documents

Query Operators – part 1

```
db.employees.find({department:{$eq:'HR'}})
```

```
db.users.find({email:{$ne:'cathy@gmail.com'}})
```

```
db.employees.find({salary:{$gt:3000}})
```

```
db.employees.find({salary:{$gte:3000}})
```

```
db.employees.find({salary:{$gte:3000,$lt:5000}})
```

```
db.employees.find({salary:{$gt:1000},department:{$eq:'HR'}})
```

```
db.employees.find({salary:{$gt:2000},department:{$in:['HR','IT']}})
```

Query Operators – part 2

```
db.employees.find({salary:{$gt:2000},department:{$nin:['HR','IT']}})  
db.employees.find({$or:[{salary:{$gt:2000}},{department:{$eq:'HR'}}]})  
db.employees.find({$and:[{salary:{$gt:2000}},{department:{$eq:'HR'}}]})  
db.employees.find({$nor:[{salary:{$gt:2000}},{department:{$eq:'HR'}}]})  
//like and but both should be false  
db.employees.find({department:{$not:{$eq:'HR'}}})  
db.users.find({email1:{$exists:false}})
```

Update Document

```
db.employees.updateOne({email:'cathy@gmail.com'},{$set:{department:'HR'}})

db.employees.updateOne(
  { email: "ria@gmail.com" },
  {
    $set:
    {
      name: "Ria K",
      email: "ria@gmail.com",
      department: "HR",
      salary: 5000,
      location: ["FL", "LA"],
      date: Date()
    }
  },
  { upsert: true }
)

db.employees.updateMany({}, { $set: { date: Date() } })
```

Delete Document

```
db.employees.deleteOne({email:'ria@gmail.com'})
```

```
db.employees.deleteMany({email:'ria@gmail.com'})
```

Update Operators(fields)

```
db.employees.updateOne({email:'cathy@gmail.com'},{$set:{email:'cathy@hotmail.com'}})
```

```
db.employees.updateMany({},{$set:{points:0}}) -- new field
```

```
db.employees.updateMany({},{$inc:{points:70}})
```

```
db.employees.updateMany({},{$rename:{points:'score'}})
```

```
db.employees.updateMany({},{$unset:{score:""}}) //deletes the field
```

Misc – skip and limit

```
db.employees.find().skip(2)
```

```
db.employees.find().skip(2).limit(1)
```

Used for pagination

Connect to local mongodb

Type mongosh –version

Type mongosh to get prompt

show dbs

use myproj to create or access new db

db.dropDatabase("dbname") to delete database (or db.dropDatabase())

show collections

db.createCollection("employees")

db.employees.drop() to delete collection

db.restaurant.renameCollection('restaurants') //rename collection

Query Operators - 3

```
db.employees.find(  
  {department:{$in:["HR","Admin"]}}  
)
```

```
db.employees.find(  
  {department:{$nin:["HR","Admin"]}}  
)
```

Update Operators (arrays)

```
db.employees.updateOne({email:'cathy@hotmail.com'},{$addToSet:{location:'F  
L'}}) //duplicates won't be added, use push instead  
  
db.employees.updateOne({email:'cathy@hotmail.com'},{$pop:{location:1}}) -try  
-1  
  
db.employees.updateMany({email:'cathy@hotmail.com'},{$pull:{points:{$gt:1}}}  
)  
  
db.employees.updateMany({email:'cathy@hotmail.com'},{$push:{points:5}})
```

Indexes (improves search but slows insert, update)

```
db.users.find({email:'cathy@gmail.com'}).explain("executionStats")
totalDocsExamined: 13,
```

```
db.users.createIndex({email:1}) //ascending
totalDocsExamined: 3,
```

```
db.users.getIndexes()
```

```
db.users.createIndex({'email':1},{unique:true})
```

```
db.users.dropIndex("email_1")
```

Aggregation pipeline

```
db.employees.aggregate([  
    {pipeline1 or stage 1 },  
    {pipeline2 or stage 2},  
])
```

Aggregation - \$match

```
db.employees.aggregate([
  {
    $match: {} //stage 1
  },
  {
    $group: { _id: "$department", total: { $sum: "$salary" } } //stage 2
  },
  {
    $sort: { "department": -1 }
  },
])
```

Sorting - collation

Product.find()

```
.collation({ locale: 'en', strength: 2 }) // strength: 2 = case-insensitive  
.sort({ name: 1 });
```

Aggregation - \$sort

```
db.employees.aggregate([
  {
    $sort: { "name": -1 }
  },
  {
    $project: {
      "name": 1,
      "email": 1,
      "salary":1
    }
  },
  {
    $limit: 5
  }
])
```

Aggregation - \$match

```
db.employees.aggregate([
  {
    $match: { salary: { $gt: 1000 } } //state 1
  },
  {
    $group: { _id: "$department", total: { $sum: "$salary" } } //stage 2
  }
])
```

Aggregation - \$group

An aggregation pipeline return results for groups of documents. For example, return the total, average, maximum, and minimum values.

```
db.employees.aggregate([
  {
    $group: {
      _id: "$department",
      Total: { $sum: "$salary" },
      Highest: { $max: "$salary" },
      Lowest: { $min: "$salary" },
      Average: { $avg: "$salary" },
    },
  },
]);
```

Aggregation - \$limit

```
db.employees.aggregate([
  { $group: { _id: "$department", Total: { $sum: "$salary" } } },
  { $limit: 1 },
]);
```

Aggregation - \$project

```
db.employees.aggregate([
  {
    $project: {
      "name": 1,
      "email": 1,
      "salary": 1
    }
  },
  {
    $limit: 2
  }
])
```

\$project – remove field

```
db.employees.aggregate([{ $project: { _id: 0, name: 0 } }]);
```

\$project – rename & add calc

```
db.employees.aggregate([
  {
    $project: {
      empname: "$name",
      email: 1,
      salary: 1,
      AnnualSalary: { $multiply: [12, "$salary"] },
    },
  },
]);
```

\$rename field (object)

```
db.profile.insertOne({  
  _id:ObjectId("6980326a71d308ebb027caa2"),  
  address:{line1:"3356 Lane 1",city:"Columbus",State:"OH"},  
  phone:7564555  
})  
  
db.profile.find(  
  {},  
  {address1:"$address.line1",city:"$address.city"}  
)
```

\$unwind – Convert Array to Object

```
db.empSkills.insertOne({  
  skills:["Java","Python",".NET"]  
})  
  
db.empSkills.aggregate([  
  {$unwind:"$skills"}  
])
```

Aggregation -\$cond

```
{ $cond: [ <boolean-expression>, <true-case>, <false-case> ] }
```

```
.....
```

```
db.employees.aggregate([
  {
    $project: {
      _id: 0,
      name: 1,
      salary: 1,
      grade: { $cond: [{ $gte: ["$salary", 2000] }, "Grade A", "Grade B"] },
    },
  },
]);
```

Aggregation - \$cond-if

```
{ $cond: { if: <boolean-expression>, then: <true-case>, else: <false-case> } }
```

```
.....
```

```
db.employees.aggregate([
  {
    $project: {
      _id: 0,
      name: 1,
      salary: 1,
      grade: {
        $cond: {
          if: { $gte: ["$salary", 2000] },
          then: "Grade A",
          else: "Grade B",
        },
      },
    },
  },
]);
```

Switch case - syntax

```
Grade: {  
  $switch: {  
    branches: [  
      { case: <boolean-expression>, then: <result> },  
      { case: <boolean-expression>, then: <result> }  
      // more branches...  
    ],  
    default: <result>  
  }  
}
```

Switch case

```
db.users.aggregate([
  {
    $project: {
      name: 1,
      level: {
        $switch: {
          branches: [
            { case: { $gte: ["$score", 90] }, then: "A" },
            { case: { $gte: ["$score", 75] }, then: "B" },
            { case: { $gte: ["$score", 60] }, then: "C" }
          ],
          default: "Fail"
        }
      }
    }
  }
])
```

Aggregation - \$lookup prep

```
db.createCollection("orders")
```

```
db.orders.insertOne({'empid':ObjectId('65fc6dd2198f1b870853d26e'), 'date':  
Date(), 'orderValue':5000})
```

Aggregation - \$lookup

```
db.orders.aggregate([
  {
    $lookup: {
      from: "employees",
      localField: "empid",
      foreignField: "_id",
      as: "employee_details",
    },
  },
])
```

Aggregation - \$lookup – pipeline

```
db.orders.aggregate([
  {$lookup: {
    from:"employees",
    let:{uid:"$empid"},
    pipeline:[
      {$match:{$expr:{$eq:[ "$_id", "$$uid" ]}}}
    ],
    as:"users"
  }}
])
```

Aggregation - \$lookup – pipeline - project

```
db.orders.aggregate([
  {
    $lookup: {
      from: "employees",
      let: { uid: "$empid" },
      pipeline: [
        { $match: { $expr: { $eq: ["$_id", "$$uid"] } } },
        {
          $project: {
            _id: 0,
            name: 1,
          },
        },
      ],
      as: "users",
    },
    { $unwind: "$users" },
    {$project:{
      name:"$users.name",
      orderValue:1
    }}
  });
})
```

Aggregation - \$out (creates ratingbydep collection)

```
db.employees.aggregate([
  {
    $project: {
      name: 1,
      department: 1,
      rating:{$convert:{input:"$rating",to:"int"}}
    },
  },
  { $group: { _id: "$department", avg: { $avg: "$rating" } } },
  {$out:"ratingByDep"}
]);
```

Views

```
db.createView(  
  "activeUsers",  
  "users",  
  [  
    { $match: { isActive: true } },  
  ]  
)  
db.activeUsers.find()  
db.activeUsers.drop()
```

LMS case study - 1

```
//use lms  
  
db.users.insertOne({  
    _id: "u1",  
    name: "Rahul Dev",  
    email: "rahul@gmail.com",  
    password: "1234",  
    role: "student", // student | instructor | admin  
});
```

LMS case study - 2

```
db.users.insertOne({  
  _id: "u2",  
  name: "Aryan",  
  email: "aryan@gmail.com",  
  password: "1234",  
  role: "instructor", // student | instructor | admin  
});
```

LMS case study - 3

```
db.users.insertOne({  
  _id: "u3",  
  name: "admin",  
  email: "admin@gmail.com",  
  password: "1234",  
  role: "admin", // student | instructor | admin  
});
```

LMS case study - 4

```
//course collection  
db.courses.insertOne({  
    _id: "c1",  
    title: "MongoDB for Beginners",  
    description: "Learn MongoDB from scratch",  
    instructorId: "u2",  
    price: 1999,  
});
```

LMS case study - 5

```
db.courses.insertOne({  
  _id: "c2",  
  title: "Nodejs for Beginners",  
  description: "Learn Nodejs",  
  instructorId: "u2",  
  price: 2000,  
});
```

LMS case study - 6

```
//modules collection  
db.modules.insertOne({  
    _id: "m1",  
    courseId: "c1",  
    title: "Introduction to MongoDB",  
    order: 1,  
});
```

LMS case study - 7

```
db.modules.insertOne({  
  _id: "m2",  
  courseld: "c1",  
  title: "CRUD Operation",  
  order: 2,  
});
```

LMS case study - 8

```
db.modules.insertOne({  
  _id: "m3",  
  courseId: "c1",  
  title: "Aggregate Pipelines",  
  order: 3,  
});
```

LMS case study - 9

```
//lesson collection  
db.lessons.insertOne({  
    _id: "l1",  
    moduleId: "m1",  
    title: "What is MongoDB?",  
    description: "MongoDB is a document database.",  
    order: 1,  
});
```

LMS case study - 10

```
//enrollment collection  
db.enrollments.insertOne({  
    studentId: "u1",  
    courseId: "c1",  
});
```

LMS case study - 11

```
//lesson progress collection  
db.lessonProgress.insertOne({  
    studentId: "u1",  
    lessonId: "l1",  
    isCompleted: true,  
});
```

LMS case study - 12

```
//quizzes collection  
  
db.quizzes.insertOne({  
    lessonId: "l1",  
    questions: [  
        {  
            question: "MongoDB is?",  
            options: ["SQL DB", "NoSQL DB", "File System"],  
            correctAnswer: "NoSQL DB",  
        },  
    ],  
});
```

Backup and Restore - Tool

Download MSI version using below link:

<https://www.mongodb.com/try/download/database-tools>

Click on the downloaded file and install

Setup environment variables to add path

C:\Program Files\MongoDB\Tools\100\bin

Backup Steps

//backup of a particular database

mongodump -d mydb -o d:/bck //d means data

//backup of a particular collection

mongodump -d mydb -c employees -o d:/bck //c means collection

//backup of all the databases

mongodump -o d:/bck //o means output

Restore Steps

```
//to restore a particular database  
mongorestore -d mydb d:/bck/mydb
```

```
//to restore a particular collection  
mongorestore -d mydb -c employees d:\bck\mydb\employees.bson
```

```
//to restore all the databases  
mongorestore --dir d:\bck\
```

```
//creates a new database and then restores  
mongorestore -d mydbnew -c employees d:\bck\mydb\employees.bson
```

```
//creates a new collection and then restores  
mongorestore -d mydbnew -c employees d:\bck\mydb\employees.bson
```

Collation with aggregate – case insensitive

```
db.collection.aggregate(  
  [{ $match: { name: "apple" } }, { $sort: { name: 1 } }],  
  { collation: { locale: "en", strength: 2 } },  
);
```

Schema Validation - 1

binary Javascript Object Notation (bson)

```
db.createCollection("customers", {  
  validator: {  
    $jsonSchema: {  
      bsonType: "object",  
      required: ["name", "email", "age"],  
    },  
  },  
});
```

db.emp1.insertOne({name:3}) – will show validation error

Schema Validation - 2

```
db.createCollection("customers", {  
  validator: {  
    $jsonSchema: {  
      bsonType: "object",  
      required: ["name", "age"],  
      properties:{  
        name:{  
          bsonType:"string",  
        },  
        age:{  
          bsonType:"int",  
        }  
      }  
    }  
  }  
});
```

Schema Validation - 3

```
db.createCollection("customers", {  
  validator: {  
    $jsonSchema: {  
      bsonType: "object",  
      required: ["name", "score"],  
      properties: {  
        name: {  
          bsonType: "string",  
        },  
        score:{  
          bsonType:["double","int","null"]  
        }  
      },  
    },  
  },  
});
```

bsonType values

Type	Meaning
string	Text
int	32-bit integer
double	Decimal
bool	true / false
date	ISODate
object	Embedded document
array	List
objectId	ObjectId

MongoDB – Regex

```
db.employees.find({name:{$regex:'Cathy'}}) //consists Cathy
```

```
db.employees.find({name:{$regex:"cathy",$options:"i"}}) // case insensitive
```

```
db.employees.find({name:{$regex:^C}}) // starts with C
```

```
db.employees.find({name:{$regex:y$}}) //ends with y
```

Mongodb cluster

Replica Set

Replica of data is created

Sharded cluster

Parts of data is stored in different machine,..used in very large database

Mongodb Replication - 1

Create a folder mongo-replica and sub folders data1 data2 and data3

Open command prompt and start running servers on separate tabs

```
mongod -replSet rs1 -logpath "d:\mongo-replica\data1\1.log" --dbpath  
"d:\mongo-replica\data1" --port 27018
```

```
mongod -replSet rs1 -logpath "d:\mongo-replica\data2\2.log" --dbpath  
"d:\mongo-replica\data2" --port 27019
```

```
mongod -replSet rs1 -logpath "d:\mongo-replica\data3\3.log" --dbpath  
"d:\mongo-replica\data3" --port 27020
```

Mongodb Replication - 2

Follow these instructions to configure replica set:

```
mongosh - -port 27018
```

```
rs.initiate({_id:"rs1",members:[{_id:0,host:"127.0.0.1:27018"},{_id:1,host:"127.0.0.1:27019"},{_id:2,host:"127.0.0.1:27020"}]})
```

```
rs.config() //to check the config
```

```
rs.status()
```

Mongodb Replication - 3

Use mongosh command with the following connection string and the primary server will automatically get connected:

```
mongosh  
"mongodb://localhost:27018,localhost:27019,localhost:27020/?replicaSet=r  
s1"
```

```
show dbs
```

```
use mytestdb
```

```
db.createCollection("customers")
```

```
db.customers.insertOne({name:"John"})
```

Mongodb Replication - 4

Check secondary servers. Check both the servers if data is replicated

mongosh --port 270xx

Secondary will start, can read but cannot write

db.getMongo().setReadPref("secondary") //or rs.secondaryOk()

use mytestdb

db.customers.find() – will work now

mongosh --port 270xx

Secondary will start, can read but cannot write

db.getMongo().setReadPref("secondary") //or rs.secondaryOk()

use mytestdb

db.customers.find() – will work now

Mongodb Replication - 5

Shutdown primary server and the primary will be automatically changed to one of the other two servers

Go to primary 270xx
Use admin
db.shutdownServer()

Now go to secondary servers 270xx or 270xx, and type show dbs...you would notice that one of the servers will be changed to primary automatically

Open new tab and start previous primary 270xx again

```
mongod -replSet rs1 -logpath d:\mongo-replica\data1\1.log --dbpath d:\mongo-replica\data1\ --port 270xx
```

Open another tab and run mongosh. You will observe that it is now a secondary server.

```
mongosh --port 270xx
```

Transactions - Commit

```
mongosh
"mongodb://localhost:27018,localhost:27019,localhost:27020/hdfc?replicaSet=rs1"

db.customers.insertOne({_id:1,name:"John",bal:500})
db.customers.insertOne({_id:2,name:"Mike",bal:100})
const session = db.getMongo().startSession();
session.startTransaction()
var custCollection = session.getDatabase("hdfc").customers
custCollection.updateOne({_id:1},{$inc:{bal:-100}})
custCollection.updateOne({_id:2},{$inc:{bal:100}})
session.commitTransaction()
session.endSession()
db.customers.find()
exit
```

Transactions - Abort

```
mongosh
"mongodb://localhost:27018,localhost:27019,localhost:27020/hdfc?replicaS
et=rs1"

const session = db.getMongo().startSession();
session.startTransaction()
var custCollection = session.getDatabase("hdfc").customers
custCollection.updateOne({_id:1},{$inc:{bal:-100}})
// don't run this - custCollection.updateOne({_id:2},{$inc:{bal:100}})
session.abortTransaction()
session.endSession()
db.customers.find()
```